

# Introduction to PDC environment

Xin Li

PDC Center for High Performance Computing  
KTH Royal Institute of Technology

PDC Summer School August 2018



# Outline

- 1 PDC Overview
- 2 Infrastructure
  - Beskow
  - Tegner
- 3 Accounts
  - Time allocations
  - Authentication
- 4 Development
  - Building
  - Modules
  - Programming environments
  - Compilers
- 5 Running jobs
  - SLURM
- 6 How to get help



# History of PDC

Year	rank	procs.	peak TFlops	vendor	name
2014	32	53632	1973.7	Cray	Beskow <sup>1</sup>
2011	31	36384	305.63	Cray	Lindgren <sup>2</sup>
2010	76	11016	92.534	Cray	Lindgren <sup>3</sup>
2010	89	9800	86.024	Dell	Ekman <sup>4</sup>
2005	65	886	5.6704	Dell	Lenngren <sup>5</sup>
2003	196	180	0.6480	HP	Lucidor <sup>6</sup>
1998	60	146	0.0934	IBM	Strindberg <sup>7</sup>
1996	64	96	0.0172	IBM	Strindberg <sup>8</sup>
1994	341	256	0.0025	Thinking Machines	Bellman <sup>9</sup>

<sup>1</sup>XC40 16-core 2.3GHz

<sup>2</sup>XE6 12-core 2.1 GHz

<sup>3</sup>XT6m 12-core 2.1 GHz

<sup>4</sup>PowerEdge SC1435 Dual core Opteron 2.2GHz, Infiniband

<sup>5</sup>PowerEdge 1850 3.2 GHz, Infiniband

<sup>6</sup>Cluster Platform 6000 rx2600 Itanium2 900 MHz Cluster, Myrinet

<sup>7</sup>SP P2SC 160 MHz

<sup>8</sup>SP2/96

<sup>9</sup>CM-200/8k



# SNIC

## Swedish National Infrastructure for Computing



National **research infrastructure** that provides a **balanced and cost-efficient** set of **resources and user support** for **large scale computation and data storage** to meet the needs of researchers from all scientific disciplines and from all over Sweden (universities, university colleges, research institutes, etc).



# Broad Range of Training

**Summer School** Introduction to HPC held every year

**Specific Courses** Programming with GPGPU, Distributed and Parallel Computing and/or Cloud Computing, Software Development Tools, CodeRefinery workshops, etc

**PDC User Days** PDC Pub and Open House



# Support and System Staff

## First-line support

Provide specific assistance to PDC users related to accounts, login, allocations etc.

## System staff

System managers/administrators ensure that computing and storage resources run smoothly and securely.

## Application Experts

Hold PhD degrees in various fields and specialize in HPC. Assist researchers in optimizing, scaling and enhancing scientific codes for current and next generation supercomputers.

# Outline

- 1 PDC Overview
- 2 Infrastructure
  - Beskow
  - Tegner
- 3 Accounts
  - Time allocations
  - Authentication
- 4 Development
  - Building
  - Modules
  - Programming environments
  - Compilers
- 5 Running jobs
  - SLURM
- 6 How to get help




# Beskow - Cray XC40 system



## Fastest machine in Scandinavia

- Lifetime: Q4 2019
- 11 racks, 2060 nodes
- Intel Haswell processor 2.3 GHz  
Intel Broadwell processor 2.1 GHz
- 67,456 cores - 32(36) cores/node
- Aries Dragonfly network topology
- 156.4 TB memory - 64(128) GB/node



- 1 XC compute blade
- 1 Aries Network Chip (4 NICs)
- 4 Dual-socket Xeon nodes
- 4 Memory DIMM / Xeon node





# Tegner

pre/post processing for Beskow

## 5 × 2TB Fat nodes

4 × 12 core Ivy Bridge, 2TB RAM  
2 × Nvidia Quadro K420

## 5 × 1TB Fat nodes

4 × 12 core Ivy Bridge, 1TB RAM  
2 × Nvidia Quadro K420

## 46 Thin Nodes

2 × 12 core Haswell, 512GB RAM  
Nvidia Quadro K420 GPU

## 9 K80 Nodes

2 × 12 core Haswell, 512GB RAM  
Nvidia Tesla K80 GPU



- Used for pre/post processing data
- Has large RAM nodes
- Has nodes with GPUs
- Has two transfer nodes
- Lifetime: Q4 2019

# Summary of PDC resources

	Beskow	Tegner
Cores in each node	32/36	48/24
Nodes	1676 Haswell 384 Broadwell	55 x 24 Haswell/GPU 10 x 48 Ivy bridge
RAM (GB)	1676 x 64GB 384 x 128GB	55 x 512GB 5 x 1TB 5 x 2TB
Allocations (core hours per month)		
Small	< 5k	< 5k
Medium	< 200k	< 80k
Large	≥ 200k	
Allocation via SNIC	yes	yes
AFS	login node only	yes
Lustre	yes	yes

# File Systems

## Andrew File System (AFS)

- Distributed file system accessible to any running AFS client
- Home directory  
`/afs/pdc.kth.se/home/[initial]/[username]`
- Access via Kerberos tickets and AFS tokens
- **Not accessible to compute nodes on Beskow**

## Lustre File System (Klemming)

- Open-source massively parallel distributed file system
- Very high performance (5PB storage - 130GB/s bandwidth)
- NO backup (always move data when done) NO personal quota
- Home directory  
`/cfs/klemming/nobackup/[initial]/[username]`

# Outline

- 1 PDC Overview
- 2 Infrastructure
  - Beskow
  - Tegner
- 3 Accounts
  - Time allocations
  - Authentication
- 4 Development
  - Building
  - Modules
  - Programming environments
  - Compilers
- 5 Running jobs
  - SLURM
- 6 How to get help



# Access requirements

User account either SUPR or PDC

Time allocation set the access limits

## Apply for PDC account via SUPR

- <http://supr.snic.se>
- SNIC database of persons, projects, project proposals and more
- Apply and link SUPR account to PDC
- Valid post address for password

## Apply for PDC account via PDC

- <https://www.pdc.kth.se/support> → "Getting Access"
- Electronic copy of your passport
- Valid post address for password
- Membership of specific time allocation

# Time Allocations

## Small allocation

- Applicant can be a PhD student or more senior
- Evaluated on a technical level only
- Limits is usually 5K corehours each month

## Medium allocation

- Applicant must be a senior scientist in Swedish academia
- Evaluated on a technical level only
- On large clusters: 200K corehours per month

## Large allocation

- Applicant must be a senior scientist in Swedish academia
- Need evidence of successful work at a medium level
- Evaluated on a technical and scientific level
- Proposal evaluated by SNAC twice a year

## Using resources

- All resources are free of charge for Swedish academia
- Acknowledgement **are** taken into consideration when applying
- Please acknowledge SNIC/PDC when using these resources:

### Acknowledge SNIC/PDC

The computations/simulations/[SIMILAR] were performed on resources provided by the Swedish National Infrastructure for Computing (SNIC) at [CENTERNAME (CENTER-ACRONYM)]

### Acknowledge people

NN at [CENTER-ACRONYME] is acknowledged for assistance concerning technical and implementation aspects [OR SIMILAR] in making the code run on the [OR SIMILAR] [CENTER-ACRONYM] resources.

# Authentication

## Kerberos Authentication Protocol

### Ticket

- Proof of users identity
- Users use passwords to obtain tickets
- Tickets are cached on the user's computer for a specified duration
- Tickets **should be created on your local computer**
- No passwords are required during the ticket's lifetime

### Realm

Sets boundaries within which an authentication server has authority (NADA.KTH.SE)

### Principal

Refers to the entries in the authentication server database (username@NADA.KTH.SE)



# Kerberos commands

## Normal commands:

`kinit` generates ticket

`klist` lists kerberos tickets

`kdestroy` destroys ticket file

`kpasswd` changes password

## On KTH-Ubuntu machines:

`pdcc-kinit`

`pdcc-klist`

`pdcc-kdestroy`

`pdcc-kpasswd`

```
$ kinit --forwardable username@NADA.KTH.SE
```

```
$ klist -Tf
```

```
Credentials cache : FILE:/tmp/krb5cc_500
```

```
Principal: username@NADA.KTH.SE
```

```
Issued      Expires      Flags Principal
```

```
Mar 25 09:45 Mar 25 19:45 FI krbtgt/NADA.KTH.SE@NADA.KTH.SE
```

```
Mar 25 09:45 Mar 25 19:45 FA afs/pdc.kth.se@NADA.KTH.SE
```

# Login using Kerberos tickets

Get a 7 days forwardable ticket on your local system

```
$ kinit -f -l 7d username@NADA.KTH.SE
```

Forward your ticket via ssh and login

```
$ ssh  
-o GSSAPIDelegateCredential=yes  
-o GSSAPIAuthentication=yes  
-o GSSAPIKeyExchange=yes  
username@clustername.pdc.kth.se
```

OR, when using ~/.ssh/config

```
$ ssh username@clustername.pdc.kth.se
```

Always create a kerberos ticket on your local system

<https://www.pdc.kth.se/support/documents/login/login.html>



# Outline

- 1 PDC Overview
- 2 Infrastructure
  - Beskow
  - Tegner
- 3 Accounts
  - Time allocations
  - Authentication
- 4 Development
  - Building
  - Modules
  - Programming environments
  - Compilers
- 5 Running jobs
  - SLURM
- 6 How to get help



# Compiling, Linking and Running Applications

on HPC clusters

**source code** C / C++ / Fortran ( .c, .cpp, .f90, .h )

**compile** Cray/Intel/GNU compilers

**assemble** into machine code (object files: .o, .obj )

**link** Static Libraries ( .lib, .a )

Shared Library ( .dll, .so )

Executables ( .exe, .x )

**request allocation** submit job request to SLURM queuing system

salloc/sbatch

**run** application on scheduled resources

aprun/mpirun



# Modules

The *modules package* allow for dynamic add/remove of installed software packages to the running environment

## Loading modules

```
module load  <software_name>
module add  <software_name>
module use  <software_name>
```

## Swapping modules

```
module swap  <software_name_1> <software_name_2>
```

## Unloading modules

```
module unload <software_name>
```

# Modules

## Displaying modules

### \$ module list

Currently Loaded Modulefiles:

- 1) modules/3.2.6.7
- ...
- 20) PrgEnv-cray/5.2.56

### \$ module avail *[software\_name]*

```
----- /opt/modulefiles -----  
gcc/4.8.1      gcc/4.9.1(default)  gcc/4.9.2      gcc/4.9.3      gcc/5.1.0
```

### \$ module show *software\_name*

```
----- /opt/modulefiles/gcc/4.9.1 -----  
conflict  gcc  
prepend-path  PATH /opt/gcc/4.9.1/bin  
prepend-path  MANPATH /opt/gcc/4.9.1/snos/share/man  
prepend-path  LD_LIBRARY_PATH /opt/gcc/4.9.1/snos/lib64  
setenv  GCC_PATH /opt/gcc/4.9.1  
-----
```

# Programming Environment Modules

specific to **Beskow**

```
Cray $ module load PrgEnv-cray
Intel $ module load PrgEnv-intel
GNU $ module load PrgEnv-gnu
```

```
$ cc source.c
$ CC source.cpp
$ ftn source.F90
```

Compiler wrappers : **cc CC ftn**

## Advantages

Compiler wrappers will automatically

- link to BLAS, LAPACK, BLACS, SCALAPACK, FFTW
- use MPI wrappers

## Disadvantage

Sometimes you need to edit Makefiles which are not designed for Cray

# Compiling serial and/or parallel code

specific to Tegner

## GNU Compiler Collection (gcc)

```
$ module load gcc openmpi
$ gcc -fopenmp source.c
$ g++ -fopenmp source.cpp
$ gfortran -fopenmp source.F90
$ mpicc -fopenmp source.c
$ mpicxx -fopenmp source.cpp
$ mpif90 -fopenmp source.F90
```

## Intel compilers (i-compilers)

```
$ module load i-compilers
$ icc -openmp source.c
$ icpc -openmp source.cpp
$ ifort -openmp source.F90
$ module add i-compilers intelmpi
$ mpiicc -openmp source.c
$ mpiicpcp -openmp source.cpp
$ mpiifort -openmp source.F90
```

## Portland Group Compilers (pgi)

```
$ module load pgi
$ pgcc -mp source.c
$ pgcpp -mp source.cpp
$ pgf90 -mp source.F90
```

## CUDA compilers (cuda)

```
$ module load cuda
$ nvcc source.cu
$ nvcc -arch=sm_37 source.cu
```



# Outline

- 1 PDC Overview
- 2 Infrastructure
  - Beskow
  - Tegner
- 3 Accounts
  - Time allocations
  - Authentication
- 4 Development
  - Building
  - Modules
  - Programming environments
  - Compilers
- 5 Running jobs
  - SLURM
- 6 How to get help



# How to run programs

- After login we are on a *login node* used only for:
  - submitting jobs,
  - editing files,
  - compiling small programs,
  - other computationally light tasks.
- **Never run calculations interactively on the login node**
- Instead, request compute resources *interactively* or via *batch script*
- All jobs must be connected to a time allocation
- For courses, PDC sets up a *reservation* for resources
- To manage the workload on the clusters, PDC uses a queueing/batch system



# SLURM workload manager

## Simple Linux Utility for Resource Management

- Open source, fault-tolerant, and highly scalable cluster management and job scheduling system
  - **Allocates** exclusive and/or non-exclusive access to **resources** for some duration of time
  - Provides a framework for **starting**, **executing**, and **monitoring** work on the set of allocated nodes
  - **Arbitrates contention** for resources by managing a queue
- Job Priority computed based on
  - Age** the length of time a job has been waiting
  - Fair-share** the difference between the portion of the computing resource that has been promised and the amount of resources that has been consumed
  - Job size** the number of nodes or CPUs a job is allocated
  - Partition** a factor associated with each node partition



# Interactive session

## salloc

### Request an interactive allocation of resources

```
$ salloc -A <account> -t <d-hh:mm:ss> -N <nodes>  
salloc: Granted job allocation 123456
```

### Run application on **Beskow**

```
$ aprun -n <PEs> -d <depth> -N <PEs_per_node> ./binary.x  
#PEs    - number of processing elements  
#depth  - number of threads (depth) per PE  
#PEs_per_node - PEs per node
```

### Run application on **Tegner**

```
$ mpirun -np <cores> ./binary.x
```

# Launch batch jobs

## sbatch

### Submit the job to SLURM queue

```
$ sbatch <script>  
Submitted batch job 958287
```

The script should contain all necessary data to identify the account and requested resources

### Example of request to run myexe for 1 hour on 4 nodes

```
#!/bin/bash -l  
  
#SBATCH -A summer-2017  
#SBATCH -J myjob  
#SBATCH -t 1:00:00  
#SBATCH --nodes=4  
#SBATCH --ntasks-per-node=32  
#SBATCH -e error_file.e  
#SBATCH -o output_file.o  
  
aprun -n 128 ./myexe > my_output_file 2>&1
```

# Monitoring and/or cancelling running jobs

## squeue -u \$USER

Displays all queue and/or running jobs that belong to the user

```
cira@beskow-login2:~> squeue -u cira
```

JOBID	USER	ACCOUNT	NAME	ST	REASON	START_TIME	TIME	TIME_LEFT	NODES	CPUS
957519	cira	pd.c.staff	VASP-test	R	None	2016-08-15T08:15:24	6:09:42	17:49:18	16	1024
957757	cira	pd.c.staff	VASP-run	R	None	2016-08-15T11:14:20	3:10:46	20:48:14	128	8192

## scancel [job]

Stops a running job or removes a pending one from the queue

```
cira@beskow-login2:~> scancel 957519
```

```
salloc: Job allocation 957891 has been revoked.
```

```
cira@beskow-login2:~> squeue -u cira
```

JOBID	USER	ACCOUNT	NAME	ST	REASON	START_TIME	TIME	TIME_LEFT	NODES	CPUS
957757	cira	pd.c.staff	VASP-run	R	None	2016-08-15T11:14:20	3:10:46	20:48:14	128	8192

# Outline

- 1 PDC Overview
- 2 Infrastructure
  - Beskow
  - Tegner
- 3 Accounts
  - Time allocations
  - Authentication
- 4 Development
  - Building
  - Modules
  - Programming environments
  - Compilers
- 5 Running jobs
  - SLURM
- 6 How to get help



# How to start your project

- Proposal for a small allocation
- Develop and test your code
- Run and evaluate scaling
- Proposal for a medium (large) allocation





# PDC support

- Many questions can be answered by reading the web documentation:  
<https://www.pdc.kth.se/support>
- Preferably contact PDC support by email: [support@pdc.kth.se](mailto:support@pdc.kth.se)
  - you get a ticket number.
  - always include the ticket number in follow-ups/replies  
they look like this: [SNIC support #12345]
- Or by phone: [+46 \(0\)8 790 7800](tel:+463087907800)
- You can also make an appointment to **come and visit**.



# How to report problems

[support@pdc.kth.se](mailto:support@pdc.kth.se)

- Do not report new problems by replying to old/unrelated tickets.
- Split unrelated problems into separate email requests.
- Use a descriptive subject in your email.
- Give your PDC user name.
- Be as specific as possible.
- For problems with scripts/jobs, give an example.  
Either send the example or make it accessible to PDC support.
- Make the problem example as small/short as possible.
- Provide all necessary information to reproduce the problem.
- If you want the PDC support to inspect some files, make sure that the files are readable.
- Do not assume that PDC support personnel have admin rights to see all your files or change permissions.



# Questions...?



# Hands-on exercise

## Login

- Some configuration steps are needed to log in to PDC
- Depends on OS: <https://www.pdc.kth.se/support/documents/login/login.html>
- In short, Kerberos and SSH supporting GSSAPI key exchange must be installed
- If needed, you will receive help to connect from your own laptops

## Exercises

- You will now practice key steps in using PDC resources
- The example source code and batch scripts can be found at [/afs/pdc.kth.se/home/k/kthw/Public/pdc\\_test.tar.gz](/afs/pdc.kth.se/home/k/kthw/Public/pdc_test.tar.gz), or in the GitHub repository <https://github.com/PDC-support/introduction-to-pdc> (in the intro-course branch)

# Hands-on exercise

## Exercises

The exercises below will demonstrate:

- 1 Kerberos - creating and listing tickets
- 2 Login - via command line or PuTTY, and where you end up on PDC
- 3 AFS home, information on your projects, `/cfs/klemming`
- 4 Working with modules
- 5 Compiling code
- 6 Submitting jobs
- 7 Running interactively

## Further information

- <https://www.pdc.kth.se/support>
- <https://hpc-carpentry.github.io/hpc-intro/>
- <https://software-carpentry.org/lessons/>

# Kerberos

- Two files should be created/edited to enable login to PDC, `krb5.conf` and (linux and mac) `.ssh/config`. For Windows, SSH configuration is done in PuTTY. Further information here: <https://www.pdc.kth.se/support/documents/login/configuration.html>
- Assuming that you have now configured Kerberos and SSH, create a forwardable Kerberos ticket
- List the ticket. Can you see if it is forwardable? What does all the jargon mean?



# Kerberos

## Answer

- To create a forwardable Kerberos ticket, type:

```
$ kinit -f <username>@NADA.KTH.SE
```

- To list your Kerberos tickets, type:

```
$ klist -f
```

The output will hopefully look similar to

```
Credentials cache: FILE:/tmp/krb5cc_H26527
```

```
Principal: kthw@NADA.KTH.SE
```

Issued	Expires	Flags	Principal
Aug 3 16:41:51 2017	Aug 4 16:39:50 2017	FfA	krbtgt/NADA.KTH.SE@
Aug 3 16:41:52 2017	Aug 4 16:39:50 2017	fA	afs/pdc.kth.se@NADA

The first line is your Kerberos ticket, the second line your AFS token.

# Login

- Start by logging in to Beskow
- Where are you after login?
  - What is your current directory?
  - What's the name of the login node?
- Have a look at the currently running processes on the login node





# Login

## Answer

- On Linux/MacOS, type one of these (use PuTTY on Windows):

```
$ ssh username@beskow.pdc.kth.se
```

```
$ ssh -o GSSAPIDelegateCredentials=yes -o GSSAPIKeyExchange=yes  
-o GSSAPIAuthentication=yes username@beskow.pdc.kth.se
```

- After logging in, `pwd` shows your current directory:

```
$ pwd  
/afs/pdc.kth.se/home/u/username
```

- The `hostname` command shows the hostname of the login node:

```
$ hostname  
beskow-login2.pdc.kth.se
```

- The `top` command shows a snapshot of currently running processes:

```
$ top
```

# AFS home, listing projects, /cfs/klemming

- Check your AFS disk quota (hint: you will need the `fs` command, type `fs help` to see available subcommands)
- Go to your klemming nobackup directory. Can you run `fs` there?
- Check which allocation(s) you belong to using the `projinfo` command
  - Check all options of `projinfo` using the `-h` flag
  - How much have your allocations been used, and for how long are they valid?



## AFS home, listing projects, /cfs/klemming

### Answer

- The `fs lq` command shows the name of the AFS volume of your home directory, your disk quota and usage:

```
$ fs lq
```

- `fs lq` only works on AFS. You can type `df -h .` instead, but there's no quota on your klemming usage:

```
$ df -h .
```

Filesystem	Size	Used	Avail	Use%	Mounted on
...:/klemming	5.2P	4.4P	730T	86%	/cfs/klemming

- The `projinfo` command accesses a database that contains logs of all allocations

```
$ projinfo
```

# Working with modules

- List the modules that are currently loaded (hint: type `module help`)
- List all the modules that are available on Beskow
- List all available modules named `PrgEnv`
- Swap from the `PrgEnv-cray` to the `PrgEnv-gnu` module



# Working with modules

## Answer

- Listing all loaded modules and all available modules is done like this:

```
$ module list  
$ module avail
```

- To list all modules matching a pattern (like PrgEnv), type

```
$ module avail PrgEnv
```

- To swap programming environments (i.e. compiler environments), type

```
$ module swap PrgEnv-cray PrgEnv-gnu
```

After swapping these modules, the compiler wrappers `cc`, `cc` and `ftn` point to the GNU compilers (instead of the Cray compilers)

# Compiling code

- Copy the tarball  
`/afs/pdc.kth.se/home/k/kthw/Public/pdc_test.tar.gz` to  
your nobackup directory, and unpack it
- Now compile the MPI example code `hello_world_mpi.c`
  - Do you need to load any MPI libraries?
  - What compiler will be used when using the `cc` compiler wrapper?



# Compiling code

## Answer

- Copying and extracting the tarball to your klemming directory:

```
$ cd /cfs/klemming/nobackup/u/username # replace this!  
$ cp /afs/pdc.kth.se/home/k/kthw/Public/pdc_test.tar.gz .  
$ tar zxf pdc_test.tar.gz  
$ cd pdc_test
```

- When using the compiler wrappers CC, cc and ftn, no MPI or numerical libraries need to be loaded or linked to explicitly!

```
$ cc -o hello_world_mpi hello_world_mpi.c
```

- Which compiler did we just compile with?

```
$ cc --version  
gcc (GCC) 4.9.1 20140716 (Cray Inc.)
```

# Submitting jobs

- Open the batch script `sbatch_beskow.sh` in your favorite editor (emacs or vim)
  - Set the allocation ID to `edu18.intropdc`
  - Set that the job should run on two nodes, with 16 processes running on each node
  - Set the requested time of the job to 2 minutes
- Submit the job to the SLURM queue!
- Monitor the queue to see if your job is running
- What output did you get?





# Submitting jobs

## Answer

- The allocation, number of nodes and time are set with these flags

```
#SBATCH -A edu18.intropdc  
#SBATCH -t 0:02:00  
#SBATCH --nodes=2
```

- If you want to run 32 MPI processes in total, with 16 processes on each node, both the `-n` and `-N` flags to `aprun` must be used:

```
aprun -n 32 -N 16 ./hello_world_mpi
```

- The job is submitted and monitored like this:

```
$ sbatch sbatch_beskow.sh  
$ squeue -u <username>
```

- The output gets written to a default filename:

```
$ cat slurm-2559495.out
```

# Running interactively

- For this example, use the OpenMP code `hello_world_openmp.c`.
  - First compile it. Does it work to compile with  
`cc -o hello_world_openmp hello_world_openmp.c?`
  - Spoiler: it won't work, since you need to add an OpenMP compiler flag. Have a look in the Makefile to find the flag
- Now allocate an interactive node. Hint: you need to specify allocation, number of nodes and time.
- Run the compiled `hello_world_openmp` on the interactive node. Important: don't forget to use `aprun`
- How many OpenMP threads do you see? How do you set the number of threads?



# Running interactively

## Answer

- To compile the code with OpenMP support, do:

```
$ cc -fopenmp -o hello_world_openmp hello_world_openmp.c
```

- To allocate an interactive node for 10 minutes:

```
$ salloc -A edu18.intropdc -t 0:05:0 -N 1
```

- To run the code using 32 threads, first set this environment variable and then run with `aprun`

```
$ export OMP_NUM_THREADS=32  
$ aprun -n 1 -d 32 ./hello_world_openmp  
# the -d (depth) isn't actually needed here
```

NOTE: `aprun` sends the job to the interactive compute node.  
Omitting `aprun` means the calculation will run on the login node.

# Questions...?

